

A Sovereign Peer-to-Peer Digital Money System

Suraj Datir
surajdatir3@gmail.com
www.sahyadri.io

Abstract. A purely peer-to-peer version of electronic cash, built as a sovereign digital money system, would allow online payments to be sent directly from one party to another without going through a financial institution. Existing proof-of-work networks solve the double-spending problem but limit throughput to a single linear chain, discarding valid blocks when multiple miners discover them simultaneously. We propose a solution based on a directed acyclic graph structure, which accepts every valid block and orders them into a single deterministic history through a novel consensus protocol. The network is secured by a memory-hard proof-of-work algorithm that compresses the gap between specialized and general-purpose hardware. State is managed through an Account + Object model, enabling direct balance transfers without tracking individual unspent outputs, while a strictly capped supply with disinflationary emission ensures predictable monetary economics. Furthermore, Sahyadri natively integrates Web5 Decentralized Identifiers (DIDs) within its Object model, providing a sovereign identity layer alongside its monetary system without introducing network state bloat. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave and rejoin the network at will, accepting the finalized DAG ordering as proof of what happened while they were gone.

1. Introduction

Sahyadri Core introduces CSM — a sovereign digital money designed as a resilient, auditable, and highly accessible settlement layer. Sahyadri is a peer-to-peer monetary system built to enable trust-minimized value transfer at global scale while preserving simplicity, predictability, and broad participation. While avoiding the severe state-bloat of general-purpose Turing-complete smart contracts, the protocol focuses on secure, low-friction digital money coupled with native Web5 Decentralized Identity (DID) primitives. Sahyadri is implemented on SahyadriDAG, a directed acyclic graph (DAG) data structure that permits parallel block creation by independent miners. Unlike

traditional single-chain blockchains that discard concurrent blocks as orphans, SahyadriDAG embraces them: every valid block contributes to the ledger, and the Sahyadri Consensus protocol produces a single deterministic total ordering of all transactions.

This eliminates the wasted-hash problem that plagues single-chain designs. The network is secured by SahyadriX — an application-layer Proof-of-Work combining Blake3 (super-speed hashing) with an 8-stage XOR memory loop requiring 16 MB of random-access memory per hash operation. This architecture makes SahyadriX genuinely memory-hard, compressing the performance gap between ASICs and general-purpose hardware (GPUs, CPUs) ensuring mining remains accessible to a broad population of participants. Sahyadri operates on a hybrid Account + Object state model. Rather than tracking individual unspent outputs (UTXOs), the protocol maintains per-address balances and nonces. Each transaction atomically deducts from the sender and credits the receiver, with balance finality enforced at the block confirmation boundary. This model delivers account + object based ergonomics with the security of Proof-of-Work.

1.1 Key Properties

Property	Specification
Ticker Symbol	CSM (Cryptographic Sovereign Money)
Smallest Unit	Kana (1 CSM = 100,000,000 Kana, 8
Maximum Supply	21,000,000 CSM (hard cap, protocol-
Block Time	1 second (deterministic)
Throughput	10,000 TPS (max_block_mass: 30,000,000)
Initial Block Reward	0.08318123 CSM per block (8,318,123
Halving Interval	Every 4 years (126,230,400 blocks)
Block Reward Split	98% Miner / 2% Development Fund
TX Fee Split	90% Miner / 10% Development Fund
Minimum TX Fee	0.00001 CSM (1,000 Kana) — fixed flat
Consensus	Sahyadri Consensus
Mining Algorithm	SahyadriX (Blake3 + 16 MB memory
State Model	Account + Object (balance + nonce per
Address Format	CSM32 (Bech32-derived, prefix: csm1...)
Network Ports	gRPC: 26110 P2P: 26111 wRPC: 27110

2. Account + Object State Model

Sahyadri employs a hybrid Account + Object state model stored in RocksDB. Unlike pure UTXO systems, which track individual unspent outputs, the account model maintains a global state table where each address has an associated balance (in CSM) and a monotonically increasing nonce. This eliminates the "dust" problem inherent in high-frequency UTXO systems at 1-second block times.

2.1 Account State

Each account entry in RocksDB contains:

- address: CSM32-encoded public key hash (csm1...)
- balance: Current spendable balance in CSM (8 decimal precision)
- nonce: Transaction sequence number (prevents replay attacks)

When a transaction is confirmed in a block, the state transition is:

```
sender.balance    -= (amount + fee)
receiver.balance  += amount
miner.balance     += fee * 0.90    // 90% of fee
company.balance   += fee * 0.10    // 10% of fee
sender.nonce      += 1
```

2.2 Object Model

The Object layer complements accounts by storing arbitrary state commitments in a Merkle trie rooted in each block header. Each block header commits to an object-based state root — a hash that cryptographically summarises the entire current state of the ledger. This enables:

- Efficient light-client verification via Merkle proofs
- Safe pruning: once a state root is committed, older block bodies may be discarded
- Fast sync: new nodes download a verified state snapshot rather than replaying all history
- Native Web5 Identity Anchoring: The Object layer serves as the cryptographic registry for Decentralized Identifiers (DIDs). Users can anchor their lightweight DID documents on-chain, acting as the foundational layer for Web5 Verifiable Credentials without congesting the consensus layer.

2.3 Why Not Pure UTXO?

At 1-second block times and 10,000 TPS, pure UTXO models generate millions of small outputs ("dust") that fragment user balances. The account model resolves this: each address maintains a single unified balance. Atomic state transitions at block confirmation boundaries ensure that partial updates are impossible — either the full transaction succeeds or the state remains unchanged.

3. Transactions

A Sahyadri transaction is a cryptographically signed state-transition instruction. It authorizes a transfer of value from a sender account to a receiver account, subject to nonce validation, balance sufficiency, and fee payment. Transactions are pending until included in a miner-confirmed block; only then do balance changes take effect.

3.1 Transaction Structure

```
Transaction {
  sender:   CSM32 address           // csmlq...
  receiver: CSM32 address           // csmlq...
  amount:   u64 (in Kana)           // 1 CSM = 100,000,000 Kana
  fee:      u64 (in Kana)           // fixed: 1,000 Kana = 0.00001 CSM
  nonce:    u64                     // must equal sender.current_nonce
  tx_id:    SHA256(sender+receiver+amount+nonce+timestamp)
}
```

3.2 Full Transaction Lifecycle

The following diagram describes the complete lifecycle from user initiation to finalized state:

```
USER INITIATES TRANSACTION
|
+--> wallet_api: POST /api/send-csm
|
+--> [VALIDATION]
|   |-- Sender account exists?
|   |-- balance >= (amount + 0.00001 CSM fee)?
|   |-- nonce == sender.current_nonce?
|   |-- No pending TX would overdraft balance?
|
+--> [MEMPOOL INSERT]
|   |-- pending_transactions table (status='pending')
|   |-- transactions table          (status='pending')
|   |-- sender.nonce += 1 (double-send prevention)
|
+--> Return: { txid, status:'pending', fee:0.00001 }
|
| (waiting for miner...)
|
MINER MINES NEXT BLOCK (1 second)
|
+--> indexer_grpc.py: flush_second()
|
+--> [BLOCK REWARD]
|   |-- payload decoded -> actual_reward_kana
|   |-- miner.balance += reward * 0.98
|   |-- company.balance += reward * 0.02
|
+--> [PENDING TX CONFIRM] (up to 10,000 per second)
|   |-- sender.balance -= (amount + fee) [NOW deducted]
|   |-- receiver.balance += amount
|   |-- miner.balance += fee * 0.90
|   |-- company.balance += fee * 0.10
|   |-- tx status = 'confirmed', block_hash = current
|   |-- If sender.balance < (amount+fee): REJECT TX
```

3.3 Mempool and Pending Queue

The mempool is implemented as a PostgreSQL table (pending_transactions) rather than an in-memory structure. This design provides:

- Persistence: pending TXs survive node restarts
- Atomic operations: PostgreSQL ON CONFLICT DO NOTHING prevents duplicate inserts
- Double-spend protection: pending locked amount is checked before accepting new TXs
- Capacity: 10,000 TX processed per 1-second block (LIMIT 10000 in confirm query)

Double-spend prevention logic in wallet_api.py:

```
# Check total locked in pending queue
SELECT COALESCE(SUM(amount + fee), 0)
FROM pending_transactions
WHERE from_address = %s AND status = 'pending'

# New TX accepted only if:
sender.balance >= pending_locked + (new_amount + fee)
```

3.4 Mempool Fairness & Anti-Bot Shield

- Per-Account Rate Limit: Maximum 50 pending transactions per account in the mempool.
- Zero-Pending VIP Rule (Network Emergency): If mempool reaches 80% capacity, only addresses with 0 pending transactions are allowed to submit new TXs. This guarantees 100% uptime for genuine users during spam attacks.
- API Spam Shield: Wallet API enforces the 50 TX limit and returns HTTP 429 (Too Many Requests) to prevent database flooding before it hits the Rust node.

3.5 Nonce and Replay Attack Prevention

Every account carries a monotonically increasing nonce. The wallet_api.py validates that the submitted nonce equals the current on-chain nonce. Each accepted TX increments the nonce immediately (before block confirmation), preventing replay of the same transaction:

```
# wallet_api.py: nonce check
if nonce != current_nonce:
    return error('Invalid nonce. Expected: %d' % current_nonce)

# On mempool accept:
UPDATE accounts SET nonce = nonce + 1 WHERE address = %s
```

3.6 Fee Model

Sahyadri uses a fixed flat fee of 0.00001 CSM (1,000 Kana) per transaction, regardless of amount. This is the lowest transaction fee of any major blockchain:

Method	Fee Model	Fee (for \$100 transfer)
Sahyadri	Fixed flat (Kana)	0.00001CSM (constant & negligible)
Payment Application	Percentage-based	\$1.50 - \$3.00
Bank Transfer (domestic)	Flat or percentage	\$1.00 - \$10.00
Bank Wire (international)	Flat + hidden fees	\$15.00 - \$50.00
Money Transfer Service	Flat + percentage	\$5.00 - \$30.00

Even at CSM = \$10,000 USD, the fee is only \$0.10 per transaction. The fee is defined once in code and applies universally:

```
# wallet_api.py + indexer_grpc.py
TX_FEE_KANA      = 1000      # 0.00001 CSM - fixed forever
TX_FEE_CSM       = 1000 / 1e8
TX_FEE_MINER_SPLIT = 0.90    # 90% to miner
TX_FEE_COMPANY_SPLIT = 0.10  # 10% to development fund
```

4. SahyadriX — Proof of Work

SahyadriX is the application-layer Proof-of-Work function used by the Sahyadri network. It combines Blake3 (a cryptographic hash function offering SIMD-optimized speed) with an 8-stage XOR memory loop operating over a 16 MB random-access memory pad. This construction is genuinely memory-hard: the evaluating device must maintain a 16 MB working set throughout computation, making it resistant to ASIC implementations that rely on small, fast on-chip caches.

4.1 Algorithm Design

```
// SahyadriX: Blake3 + 16MB RandomX-style Memory Loop
// File: crypto/hashes/src/pow_hashers.rs

const MEM_SIZE: usize = 16 * 1024 * 1024; // 16 MiB
const ROUNDS:  usize = 1024;              // random-access rounds
const WINDOW:   usize = 32;               // bytes per access

pub fn hash(data: &[u8]) -> Hash {
    // Step 1: Initial Blake3 hash
    let mut hasher = blake3::Hasher::new();
    hasher.update(data);
    let mut current_hash = *hasher.finalize().as_bytes();

    MEM_PAD.with(|pad| {
        let mut mem_pad = pad.borrow_mut();

```

```

// Step A: Fill 16MB memory deterministically from seed hash
let mut seed_hash = current_hash;
let mut i = 0usize;
while i < MEM_SIZE {
    let mut h = blake3::Hasher::new();
    h.update(&seed_hash);
    h.update(&(i as u64).to_le_bytes());
    let out = h.finalize();
    mem_pad[i..i+WINDOW].copy_from_slice(&out.as_bytes()
[0..WINDOW]);
    seed_hash = *out.as_bytes();
    i += WINDOW;
}

// Step B: 1024 random-access mixing rounds (ASIC-killer)
for _ in 0..ROUNDS {
    let idx1 = (u32::from_le_bytes(current_hash[0..4]
        .try_into().unwrap()) as usize) % max_index;
    let idx2 = (u32::from_le_bytes(current_hash[4..8]
        .try_into().unwrap()) as usize) % max_index;
    let mut h = blake3::Hasher::new();
    h.update(&current_hash);
    h.update(&mem_pad[idx1..idx1+WINDOW]);
    h.update(&mem_pad[idx2..idx2+WINDOW]);
    current_hash = *h.finalize().as_bytes();
}
});
Hash::from_bytes(current_hash)
}

```

4.2 Device Balance

SahyadriX is calibrated so that the performance differential between dedicated mining ASICs and general-purpose hardware is compressed. The 16 MB memory requirement eliminates the ASIC advantage from cache size. CPU & GPU can all participate competitively in Sahyadri mining.

Device Class	Memory Available	SahyadriX	Relative Efficiency
ASIC (custom)	Limited on-chip	Constrained by 16 MB req.	~2-3x GPU
GPU (consumer)	4+ GB VRAM	Full parallel	Baseline
CPU (modern)	System RAM	Single-thread	~0.3-0.5x GPU

4.3 Ignition Switch — Connecting PoW to Block Template

```

// PowHash: binds PoW to specific block contents
pub fn finalize_with_nonce(&self, nonce: u64) -> Hash {
    let mut data = Vec::with_capacity(48);
    data.extend_from_slice(self.pre_pow_hash.as_bytes()); // block
contents
    data.extend_from_slice(&self.timestamp.to_le_bytes());
    data.extend_from_slice(&nonce.to_le_bytes());
    SahyadriX::hash(&data) // 16MB computation over this exact
template
}

```

5. Network Architecture


The Sahyadri network operates as a fully decentralized peer-to-peer system. Nodes communicate over three primary channels:


- gRPC (port 26110): Client-node communication; used by the indexer and wallet API
- P2P (port 26111): Block and transaction propagation between nodes worldwide
- wRPC / WebSocket (port 27110): Browser-based and light-client interfaces

5.1 Node Types

Node Type	Data Retained	Use Case	Storage
Full Archive Node	All blocks + full TX history + current state	Explorer, auditing	Growing (GB-TB)
Pruned Full Node	Current state + recent blocks (RocksDB)	Mining, validation	5-10 GB (stable)
Light Client	Block headers + Merkle proofs only	Wallet verification	Minimal

5.2 Data Flow Architecture






6. Consensus — SahyadriDAG

Sahyadri Consensus is a deterministic finality engine that operates on the SahyadriDAG data structure to produce a single, total, immutable ordering of all blocks and transactions. Unlike probabilistic longest-chain consensus, Sahyadri provides absolute finality: once a block is finalized, its position in the ordering is permanent.

6.1 Parallel Block Production

Because SahyadriDAG permits concurrent block proposals, multiple miners may simultaneously mine different blocks at the same height. All valid blocks contribute to the ledger — none are discarded as orphans. The Sahyadri Consensus algorithm traverses the DAG to compute a deterministic total ordering:



The indexer handles parallel blocks with a second-buffer mechanism: all blocks arriving within the same 1-second window are grouped, and a single combined reward is issued to the miner of the representative block (best_block_hash):

```

# indexer_grpc
second_buffer = defaultdict(list) # timestamp_sec -> [blocks]

def process_block(block, conn):
    timestamp_sec = timestamp_ms // 1000
    second_buffer[timestamp_sec].append(block)
    if timestamp_sec > last_second:
        flush_second(last_second, second_buffer[last_second], conn)
  
```

6.2 Finality Guarantees

Sahyadri Consensus provides the following guarantees:

- Deterministic: Every honest node computing the same DAG arrives at identical ordering
- Irreversible: Finalized blocks cannot be reorganized or replaced
- No orphans: All valid blocks contribute to the DAG and ledger
- Sub-second propagation: New blocks reach the global network in under 1 second

6.3 Dual Finality: Transaction vs. Reward

SahyadriDAG distinguishes between user transaction finality and miner reward finality to optimize both payment speed and network security:

- Transaction Finality (~1 Second): For users transferring CSM, deterministic finality is immediate. Once a transaction is included in a block and confirmed by the BFT layer, the balance transfer is irreversible within approximately 1 second. This enables instant Point-of-Sale (POS) payments and IPV (Section 10).

- Reward Finality & Blue Score (within a few seconds): Because SahyadriDAG permits parallel block production, multiple miners may produce blocks simultaneously. While all valid blocks contribute to the DAG, block rewards are strictly issued based on Blue Score confirmation. It takes a few seconds for the DAG to calculate the deterministic total ordering and assign Blue Scores to blocks. Miners only receive the 98% block reward and 90% fee split once their block is confirmed in the Blue Set. Red blocks contribute to network throughput but do not receive coinbase rewards.

6.4 10,000 TPS Configuration

Transaction throughput is governed by `max_block_mass` in the consensus parameters. Each account-model transaction has a mass of approximately 3,000 units:

```
// consensus/core/src/config/params.rs
// Set across all 4 network configs: mainnet, testnet, simnet, devnet
max_block_mass: 30_000_000,

// Capacity calculation:
// 30,000,000 mass / 3,000 mass-per-tx = 10,000 TX per block
// 10,000 TX per block * 1 block per second = 10,000 TPS
```

6.5 Consensus Mechanics: Achieving 1-Second BFT Finality and 10,000 TPS

• Unlocking 10,000 TPS with SahyadriDAG

Traditional blockchain architectures suffer from linear bottlenecks, where only one block can be processed and appended at a time. Sahyadri fundamentally shifts this paradigm by utilizing a Directed Acyclic Graph (DAG) structure. In the Sahyadri network, multiple miners can propose blocks concurrently. Because parallel blocks are not orphaned but instead cryptographically woven into the DAG, the network's throughput increases exponentially. This parallel processing architecture allows Sahyadri to scale up to 10,000 Transactions Per Second (TPS) without compromising the decentralized, permissionless nature of Proof-of-Work.

• Bridging PoW with 1-Second BFT Finality

Typically, Proof-of-Work relies on probabilistic finality (waiting for multiple block confirmations), which results in long settlement times. Sahyadri solves this by decoupling block generation from block finalization. While the PoW mechanism drives the rapid, parallel creation of blocks within the DAG, a deterministic Byzantine Fault Tolerance (BFT) consensus layer operates in tandem to finalize them.

Once parallel blocks are proposed, the network's validation layer rapidly orders them and cryptographically attests to their position within the graph. Because BFT consensus relies on absolute voting thresholds rather than longest-chain rules, this hybrid approach delivers irreversible, deterministic finality within approximately 1 second. Users get the instant settlement of BFT with the robust security foundation of PoW.

In live network conditions, the Rust node accepts multiple parallel blocks per second (observed at ~5-10 BPS depending on hash rate), while the Indexer calculates deterministic Blue Score finality to issue rewards and state updates, perfectly maintaining the separation between rapid block creation and immutable financial finality.

7. Block Rewards and Fee Economics

Sahyadri's incentive structure combines predictable block issuance with a flat transaction fee model. All monetary parameters are fixed at genesis and cannot be altered without a consensus-breaking hard fork.

7.1 Block Reward

Each finalized block carries an initial reward of 0.08318123 CSM (8,318,123 Kana). Rewards are split at the Rust node level before writing to RocksDB:

```
// consensus/src/processes/coinbase.rs
```

```
pub fn calc_block_subsidy(&self, blue_score: u64) -> u64 {
```

```

    let base_reward: u64 = 8_318_123;          // 0.08318123 CSM
    let halving_interval: u64 = 126_230_400; // 4 years at ~1block/sec
    let halvings = blue_score / halving_interval;
    if halvings >= 64 { return 0; }
    base_reward.checked_shr(halvings as u32).unwrap_or(0)
}

// Indexer split (indexer_grpc.py)

BLOCK_MINER_SPLIT   = 0.98 # 98% to miner
BLOCK_COMPANY_SPLIT = 0.02 # 2%  to development fund

```

7.2 Transaction Fee Split

```

# indexer_grpc.py - pending TX confirmation

TX_FEE_KANA           = 1000           # 0.00001 CSM (fixed forever)
TX_FEE_MINER_SPLIT    = 0.90           # 90% to block miner
TX_FEE_COMPANY_SPLIT  = 0.10           # 10% to development fund

# Per confirmed TX:
miner.balance += fee * 0.90
company.balance += fee * 0.10

```

7.3 Halving Schedule

Epoch (k)	Block Reward (CSM)	Years	Annual Emission (CSM)	Cumulative Supply (CSM)
0 (Genesis)	0.08318123	0-4	~437,459	~437,459
1	0.04159062	4-8	~218,729	~656,188
2	0.02079531	8-12	~109,365	~765,553
3	0.01039765	12-16	~54,682	~820,235
...	...	...	...	...
63	~0	>252	~0	~21,000,000

"Halving is triggered by Blue Score, not raw block count. Since parallel blocks exist in the DAG, raw block count is higher than Blue Score. Halving occurs when Blue Score reaches 126,230,400 (representing ~4 years of chronological time), maintaining the strict 21M cap."

The total supply converges to 21,000,000 CSM asymptotically. The geometric series property ensures:

$$S(\text{inf}) = 2 \times H \times s_0 = 2 \times 126,230,400 \times 0.08318123 = 21,000,000 \text{ CSM}$$

8. Blockchain Indexer and Explorer

Because Sahyadri nodes prune old block data (retaining only current account state in RocksDB), a separate indexer service captures and permanently archives the full transaction history in PostgreSQL. This two-tier architecture keeps nodes lightweight (5-10 GB) while the explorer maintains complete historical records.

8.1 Indexer Architecture (indexer_grpc.py)

The indexer connects to the node via gRPC on port 26110 and processes blocks in real time. It performs historical sync on startup (resuming from last processed blue_score) and then switches to live block listening:

INDEXER STARTUP SEQUENCE:

```
1. Connect to sahyadri via gRPC :26110
2. Query PostgreSQL: SELECT MAX(blue_score) FROM blocks
3. Historical sync:
  loop:
    GET /getBlocksRequest(lowHash, includeTransactions=true)
    FOR each block WHERE blue_score > last_processed:
      process_block(block)
    IF no new blocks: BREAK
    low_hash = last_block.hash
4. Subscribe to live blocks:
  notifyBlockAddedRequest -> stream
  FOR each blockAddedNotification:
    GET full block with transactions
    process_block(block)
```

8.2 Database Schema

Table	Key Columns	Purpose
blocks	block_hash, blue_score,	Block registry
transactions	tx_id, from_address, to_address, amount, fee, status, block_hash	All user TXs
pending_transactions	tx_id, from_address, to_address, amount, fee, nonce, status	Mempool queue
rewards	tx_id, block_hash, miner,	Mining rewards
accounts	address, balance, nonce	Current state cache

8.3 Explorer API (server.js — port 3000)

Endpoint	Description
GET /api/blocks	Recent blocks with pagination
GET /api/transactions	All transactions with pagination
GET /api/block/:hash	Block detail + reward
GET /api/block-by-score/:score	Block by blue_score
GET /api/address/:address	Balance + TX history
GET /api/top-addresses	Rich list (top 100)
GET /api/transaction/:hash	TX detail
GET /api/stats	Network stats (supply, wallets, TXs)

8.4 Wallet API (wallet_api.py — port 5000)

Endpoint	Description
GET /api/balance/:address	Current balance from accounts table
GET /api/history/:address	User TX history (excludes mining)
POST /api/send-csm	Submit pending transaction to mempool
GET /api/nonce/:address	Current nonce for sender address
GET /api/fee	Current network fee (TX_FEE_CSM, TX_FEE_KANA)

9. Disk Space and Pruning

Efficient disk management is essential to keeping Sahyadri nodes lightweight. The protocol architecture cleanly separates what the node keeps (current state) from what the explorer keeps (full history).

9.1 Node Storage Model

The Sahyadri node (sahyadrid) uses RocksDB exclusively for current state:

Node (RocksDB) KEEPS:

```
+-----+
| Current account balances |
| Current nonces           |
| Recent blocks (consensus)|
| DAG tips (for new blocks)|
+-----+
Size: ~5-10 GB (stable)
```

Node PRUNES (auto):

```
+-----+
| Raw block bodies (>30h)  |
| Old TX history           |
| Old UTXO sets (not used) |
| Pruned state snapshots   |
+-----+
Size: grows then discarded
```

Because Sahyadri uses an Account model (not UTXO), account pruning does not apply. Zero-balance accounts may theoretically be removed, but this is not implemented — the cost is trivial (100 bytes per account).

9.2 Explorer Storage Model

The PostgreSQL explorer database grows continuously and is never pruned — it is the permanent historical archive:

Storage estimate per million transactions:

1 TX record = ~500 bytes

1M TX = ~500 MB

100M TX = ~50 GB

1B TX (future) = ~500 GB

At current load (10-100 TX/sec):

1 day = 864,000 to 8,640,000 TX = 432 MB to 4.3 GB/day

9.3 Recommended Deployment

Sahyadri's architecture is designed to run efficiently on standard, affordable hardware, ensuring broad participation without reliance on specific cloud providers. Nodes can be deployed on any standard VPS, dedicated server, or home hardware running Linux/Mac.

Component	Minimum Spec	Recommended Spec (High Load)
Node (sahyadrid)	4 CPU, 8 GB RAM, 50 GB SSD	8 CPU, 16 GB RAM, 100 GB NVMe
Indexer (Python)	1 CPU, 2 GB RAM	2 CPU, 4 GB RAM
PostgreSQL	2 CPU, 4 GB RAM, 100 GB+	4 CPU, 8 GB RAM, 500+ GB NVMe
API Services	1 CPU, 1 GB RAM	2 CPU, 2 GB RAM

Deployment Notes:

- Pruned Nodes (Miners): Require minimal storage (~5-10 GB stable) and can comfortably run on low-end VPS or standard consumer hardware.
- Archival Nodes (Explorers/Indexers): Storage requirements grow continuously with network usage. NVMe SSDs are highly recommended for PostgreSQL to handle the 10,000 TPS batch insertions efficiently. Storage can be expanded incrementally as the network scales.

9.4 State Management: 30-Hour Pruning & Archival Offloading

High-throughput blockchains typically suffer from massive state bloat, making it prohibitively expensive to run a full node. Sahyadri addresses this through a strict Blue Score-based pruning architecture, separating lightweight consensus from archival storage.

- **The Pruned Consensus Layer (Default):** By default, Sahyadri nodes operate as pruned nodes. The protocol enforces a deterministic pruning depth calculated via a mathematical lower bound combining finality depth (12 hours), merge depth, and mergeset limits. With a default pruning duration of 108,000 seconds (30 hours), a 10 BPS network retains approximately 1,080,000 blocks of DAG history in its local RocksDB. This ensures that node storage remains extremely lightweight and stable (5-10 GB), allowing CPU/GPU resources to be dedicated entirely to securing the network rather than managing historical data.
- **High-Throughput Batch Processing (10,000 TPS Engine):** To handle the theoretical maximum of 10,000 TPS without bottlenecking at the database layer, the Sahyadri Indexer utilizes a RAM-based batch insertion engine. Instead of executing individual INSERT queries for each transaction, the indexer computes double-spend and nonce validations entirely in memory, calculating balance transitions as a unified state array. This data is then pushed to PostgreSQL using **execute_values** in massive arrays, reducing database round-trips from tens of thousands per block to just 4 optimized queries. This ensures 10,000 TPS is processed smoothly without latency.
- **Archival Offloading via PostgreSQL:** To preserve the immutability and availability of the complete transaction history for Block Explorers, CEXs, and audits, Sahyadri utilizes Archival Nodes. By running the node with the —**archival** flag, the protocol bypasses all local data deletion. The PostgreSQL Indexer connects to this Archival Node, seamlessly archiving all transactions into an external database.
- **Mempool Pruning & Data Optimization:** To prevent the mempool (**pending_transactions** table) from becoming a storage sink, the indexer implements automated 24-hour pruning. Transactions that have reached a terminal state (**confirmed** or **rejected**) are purged from the pending queue after 24 hours. This constant garbage collection ensures that the mempool only retains active, pending transactions, keeping query times minimal even under sustained high network load.
- **Security against Misconfiguration:** If an operator attempts to remove the —**archival** flag from a previously archival node, the node halts with an explicit warning and requires manual confirmation, preventing accidental data deletion and ensuring the integrity of the network's historical record.

10. Instant Payment Verification (IPV)

Instant Payment Verification allows lightweight clients to confirm transaction finality without downloading the full blockchain. Because Sahyadri uses deterministic consensus rather than probabilistic longest-chain, finality is absolute: a transaction confirmed in a block cannot be reversed.

A light client verifying a payment needs only:

- The finalized block header (contains state root and block hash)
- A compact Merkle inclusion proof linking the TX to the block
- Proof that the block has been finalized by Sahyadri Consensus

Sahyadri IPV is immediate: 1 block confirmation = absolute finality. This enables:

- Point-of-sale terminals: Accept payments as fast as block confirmation (1 second)
- Merchants: Same confidence as cash settlement
- Embedded payment processors: Minimal bandwidth and compute requirements

11. Web5 and Sovereign Identity

Sahyadri extends the concept of sovereignty beyond digital money into digital identity by natively supporting Web5 protocols at the base layer. Traditional Layer-1 networks treat identity as an afterthought, often relying on centralized off-chain servers or state-heavy smart contracts. Sahyadri integrates identity directly into its Account + Object model without compromising its 10,000 TPS capacity or 30-hour pruning architecture.

11.1 Decentralized Identifiers (DIDs)

Users can cryptographically generate and control a unique Sahyadri DID (e.g., did:sahyadri:csmls...). The DID document—containing only cryptographic public keys and service endpoints—is anchored in the Sahyadri Object state. This allows passwordless authentication and true self-sovereign ownership of digital identity.

11.2 Decentralized Web Nodes (DWNs) for State Efficiency

To maintain high throughput and prevent state bloat, heavy identity metadata (such as profile information, KYC documents, or encrypted messages) is never stored on the Sahyadri blockchain. Instead, the network utilizes Decentralized Web Nodes (DWNs). The on-chain DID simply points to the user's off-chain DWN, ensuring the Sahyadri consensus node remains ultra-lightweight (5-10 GB) while users retain absolute, censorship-resistant control over their personal data.

11.3 Verifiable Credentials (VCs) and Compliance

Sahyadri's native DID support enables seamless integration with institutional platforms, payment gateways, and Centralized Exchanges (CEXs) through Verifiable Credentials (VCs). Users can cryptographically prove real-world attributes (e.g., KYC clearance or jurisdiction) to an exchange without exposing their underlying sensitive data. This provides a compliant, privacy-preserving bridge between sovereign digital money and regulatory frameworks, solving the compliance trilemma natively.

11.4 Decentralized Web Nodes (DWN) and Encrypted Data Vaults

To complement the sovereign identity layer, Sahyadri implements a **Local-First Data Architecture** utilizing Decentralized Web Nodes (DWNs). This ensures that while the blockchain handles financial finality, the user's personal data remains under their absolute cryptographic control.

1. Sovereign Storage & Encrypted Vaults:

Users can store and backup diverse data types—including images, videos, and documents (up to 100MB per object)—within their personal DWN. These files are encrypted at the edge (on the user's device) using keys derived from their unique Sahyadri DID. Access is restricted exclusively to the DID owner; any unauthorized attempt to access the data vault via a different DID results in an automatic cryptographic rejection.

2. Peer-to-Peer (P2P) Messaging Layer:

Sahyadri natively supports DID-to-DID encrypted communication. By utilizing a "Plus-Icon" simplified interface, users can initiate secure chats by simply entering the recipient's Sahyadri DID (**did:sahyadri:csm1s...**).

- **Privacy by Design:** Messages are never stored on the central Sahyadri blockchain, preventing state bloat. Instead, they are relayed between DWNs, ensuring a chats like interface real-time experience with the privacy of a sovereign network.
- **Resilience:** All communication and file metadata are backed up locally on the user's device, allowing for seamless recovery and migration across hardware without relying on centralized cloud providers.

3. State Efficiency:

By offloading heavy data to the DWN layer, Sahyadri maintains its ultra-lightweight node profile (5-10 GB) even at 10,000 TPS, solving the storage trilemma of high-throughput networks.

12. Privacy

Privacy in Sahyadri is achieved by separating value transfer from identity. All transactions are publicly verifiable on the explorer, but the protocol does not associate transfers with real-world identities. Ownership is defined exclusively by cryptographic control of private keys. Each user generates a CSM32 address derived from their public key (secp256k1 ECDSA). Users are encouraged to generate fresh key pairs for each receiving address. The Account model naturally provides less transaction graph ambiguity than UTXO systems, but the protocol does not require identity disclosure.

Sahyadri does not include built-in mixing, confidential transactions, or zero-knowledge proofs at the base layer. Users requiring stronger privacy may layer additional tools on top of the base protocol. The L1 provides robust baseline privacy: no account registration, no identity verification, no link between real-world identity and on-chain activity. However, through the Web5 identity layer, users have the power of 'Opt-in Identity.' They can choose to selectively disclose verified attributes via Verifiable Credentials (VCs) when interacting with compliant entities or exchanges, maintaining a perfect balance between base-layer pseudonymity and institutional composability.

13. Calculations and Economic Model

13.1 Emission Formula

$$s_k = s_0 \times 2^{(-k)}$$

$$I_k = H \times s_k \quad (\text{total emission in epoch } k)$$

$$S(n) = 2 \times H \times s_0 \times (1 - 2^{(-n)})$$

$$S(\text{inf}) = 2 \times H \times s_0 = 21,000,000 \text{ CSM}$$

Where $s_0 = 0.08318123$ CSM, $H = 126,230,400$ blocks (4 years at 1 BPS), k = halving epoch.

13.2 Block Revenue

$$R_k = s_k + f_{\text{avg}} \quad (\text{total per-block revenue})$$

$$\phi_k = f_{\text{avg}} / (s_k + f_{\text{avg}}) \quad (\text{fee fraction of total revenue})$$

$$\text{AnnualRevenue}_k = r_{\text{yr}} \times (s_k + f_{\text{avg}})$$

Where $r_{\text{yr}} = 31,536,000$ blocks/year (at 1 BPS), f_{avg} = average fee revenue per block.

13.3 Fee Revenue at Scale

Daily TX Volume	Daily Fee Revenue (CSM)	Miner Share (90%)	Dev Fund (10%)
10,000 TX/day	0.1 CSM	0.09 CSM	0.01 CSM
1,000,000 TX/day	10 CSM	9 CSM	1 CSM
100,000,000 TX/day	1,000 CSM	900 CSM	100 CSM
864,000,000 TX/day (10k TPS)	8,640 CSM	7,776 CSM	864 CSM

13.4 Network Security Budget

As block subsidies halve over time, transaction fees progressively constitute a larger fraction of miner revenue ($\phi_k \rightarrow 1$). At 10,000 TPS capacity, Sahyadri is positioned to generate substantial fee revenue even at low individual fees:

Full-capacity fee revenue = 10,000 TX/sec x 0.00001 CSM x 86,400 sec = 8,640 CSM/day

14. Security Model

The Sahyadri network derives security from four pillars: verifiable computation, deterministic consensus, economic incentives, and cryptographic primitives.

14.1 Formal Security Conditions

Transaction validity: $\text{Verify}(\text{tx}) = \text{TRUE}$ iff

$\text{sig_valid}(\text{tx.sig}, \text{tx.sender_pubkey})$ AND
 $\text{accounts}[\text{tx.sender}].\text{balance} \geq (\text{tx.amount} + \text{tx.fee})$ AND
 $\text{accounts}[\text{tx.sender}].\text{nonce} == \text{tx.nonce}$

Block acceptance: $\text{Accept}(\text{B})$ iff

$\text{VerifySahyadriX}(\text{B.header})$ AND
ALL tx in B : $\text{Verify}(\text{tx}) = \text{TRUE}$ AND
 $\text{ConsensusFinalized}(\text{B}) = \text{TRUE}$

Attack cost: $P_{\text{attack}} = \alpha$ (attacker hash fraction)

For $\alpha < 0.5$: attack is economically irrational

After finalization: reorganization probability = 0 (deterministic)

14.2 Double-Spend Impossibility

Double-spend attacks are prevented at two layers:

- Mempool layer (wallet_api.py): Pending locked balance check prevents overdraft before block confirmation
- Confirmation layer (indexer_grpc.py): Sender balance checked again at confirmation; insufficient balance = TX rejected
- Nonce layer: Each TX carries a unique nonce; same nonce cannot be reused after increment
- Consensus layer: Finalized blocks are irreversible; deterministic ordering eliminates reorg risk

Network Layer Protection:

Beyond state-level validations, Sahyadri secures the network perimeter against malicious infrastructure attacks. The P2P layer (port 26111) includes an in-built IP banning mechanism designed to detect and isolate transaction flooding. If a node detects an abnormal volume of malformed or spam transactions from a specific IP, it automatically blacklists the source, protecting the node from DDoS attacks at the network edge. This ensures that malicious actors cannot choke node bandwidth or consume computational resources, maintaining high availability and deterministic 1-second block propagation even under active network attacks.

15. Governance

Sahyadri adopts a minimal governance model. All core monetary parameters are permanently fixed at genesis and are not subject to modification through proposals or voting:

- Total supply: 21,000,000 CSM — immutable
- Initial block reward: 0.08318123 CSM — immutable
- Halving interval: 126,230,400 blocks (~4 years) — immutable
- Minimum TX fee: 0.00001 CSM — immutable
- Block reward split: 98/2 — immutable
- TX fee split: 90/10 — immutable

Changing any monetary parameter requires a consensus-breaking hard fork, effectively creating a new chain. The development fund (2% block reward + 10% TX fee) is allocated to ongoing protocol development, infrastructure, and ecosystem support. No central governance body, no token-weighted voting, and no on-chain governance process exists. The protocol evolves through rough consensus and running code.

16. Conclusion

Sahyadri presents a comprehensive design for a sovereign peer-to-peer digital money system. By combining a directed acyclic graph data structure for parallel block production, a memory-hard Proof-of-Work algorithm for inclusive security, an Account + Object state model for simplified value transfer, and a fixed flat fee structure for predictable costs, the protocol achieves deterministic finality in seconds with high-throughput transaction processing. The strictly capped monetary supply, transparent halving schedule, and protocol-enforced immutable parameters provide a stable, auditable, and predictable foundation for global value transfer — digital money that works reliably, fairly, and forever.